

In [49]:

```
import numpy as np
import pandas as pd
import seaborn as sns

import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')

import plotly.express as px
import plotly.graph_objects as go
```

NOTE : Execute below cells only if running on google colab for getting the dataset

In [50]:

```
!gdown --id 192h8LDFatS7r8GHI6trlegkheGKgBVdE --output CleanDF.csv
```

Downloading...

From: <https://drive.google.com/uc?id=192h8LDFatS7r8GHI6trlegkheGKgBVdE>

To: /content/CleanDF.csv

100% 747k/747k [00:00<00:00, 50.3MB/s]

In [51]:

```
df = pd.read_csv('CleanDF.csv')
df.head(3)
```

Out[51]:

	title_translated	listed_price	retail_price	units_sold	uses_ad_boosts	rating	rating_count	
0	2020 Summer Vintage Flamingo Print Pajamas Se...	16.0	14	100	0	3.76	54	
1	Women's Casual Summer Sleeveless Sexy Mini Dress	8.0	22	20000	1	3.45	6135	
2	2020 New Arrival Women Spring and Summer Beach...	8.0	43	100	0	3.57	14	

EDA

group the features by whether they are categorical, numerical, or 'other'.

In [52]:

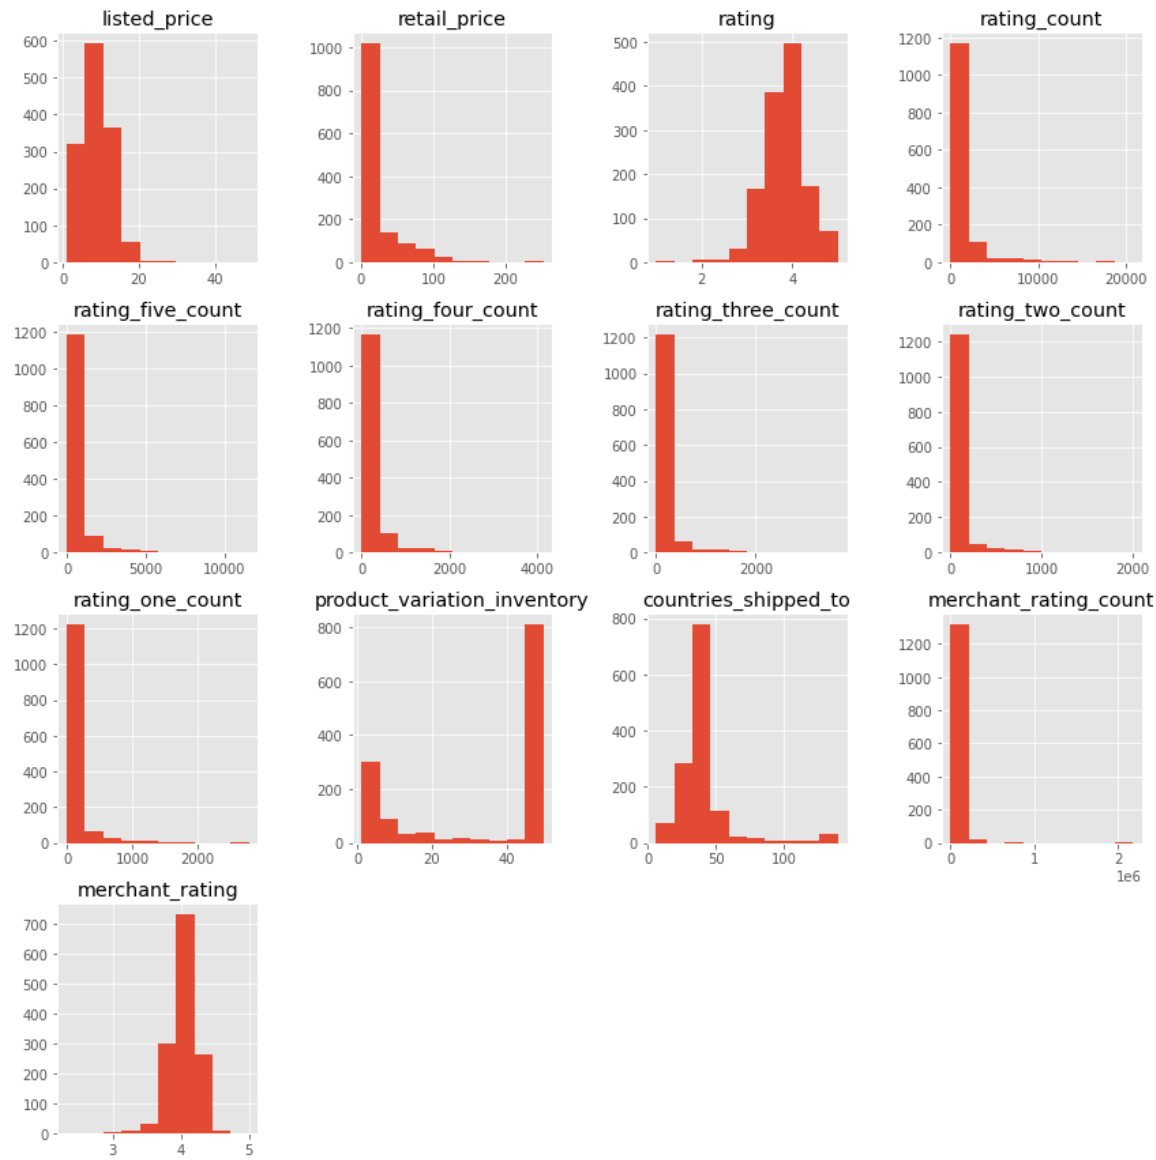
```
numerical_features = ['listed_price', 'retail_price', 'rating', 'rating_count',  
                      'rating_five_count', 'rating_four_count',  
                      'rating_three_count', 'rating_two_count', 'rating_one_count',  
                      'product_variation_inventory',  
                      'countries_shipped_to', 'merchant_rating_count', 'merchant_rating']  
  
categorical_features = ['units_sold', 'uses_ad_boosts', 'badge_product_quality',  
                        'product_color',  
                        'product_variation_size_id', 'shipping_option_price', 'has_urgency_banner',  
                        'urgency_text']  
  
other_features = ['title_orig', 'tags', 'merchant_id', 'product_picture', 'product_id']
```

Let's use `DataFrame.hist()` to visualize the distribution of the numerical features:

- The `price` and `countries_shipped_to` features are positively skewed.
- The `rating` and `merchant_rating` features are negatively skewed.
- The `retail_price`, `rating_count` ... and `merchant_rating_count` features are exponential distributions rather than Gaussian.
- The `product_variation_inventory` is organic; not obeying any probability distribution.

In [53]:

```
hist = df[numerical_features].hist(figsize=(12, 12))  
  
plt.tight_layout()  
plt.show()
```



In [54]:

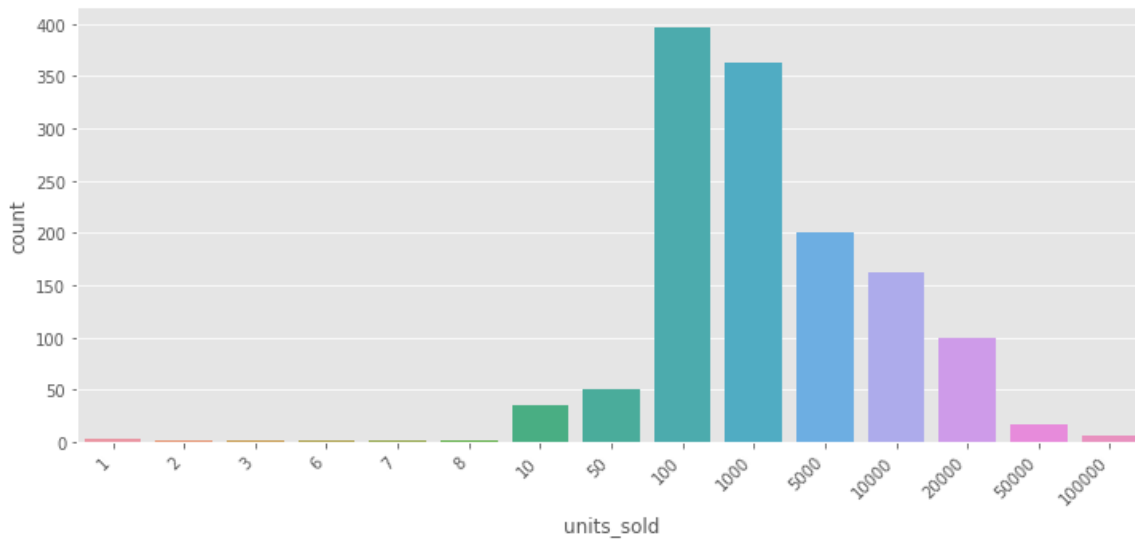
```
# visualize the categorical features.
```

```
fig, ax = plt.subplots(figsize=(12, 5))
```

```
hist = sns.countplot(x='units_sold', data=df, order=sorted(df['units_sold'].unique()), ax=ax)
```

```
plt.xticks(rotation=45, ha='right')
```

```
plt.show()
```



In [55]:

```
# The features: `uses_ad_boosts`, `badge_product_quality`, and `has_urgency_banner` are binary.

fig = plt.figure(figsize=(12, 3))

boolean_features = ['uses_ad_boosts', 'badge_product_quality', 'has_urgency_banner']
for i in range(3):
    true_percentage = (df[boolean_features[i]].value_counts()[1] / len(df[boolean_features[i]])) * 100
    print("Percent that the '%s' flag is True: %f" % (boolean_features[i], true_percentage))

    fig.add_subplot(1, 3, i + 1)
    sns.countplot(df[boolean_features[i]])

plt.tight_layout()
plt.show()
```

Percent that the 'uses_ad_boosts' flag is True: 43.549590

Percent that the 'badge_product_quality' flag is True: 7.755406

Percent that the 'has_urgency_banner' flag is True: 27.293065

/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:


Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

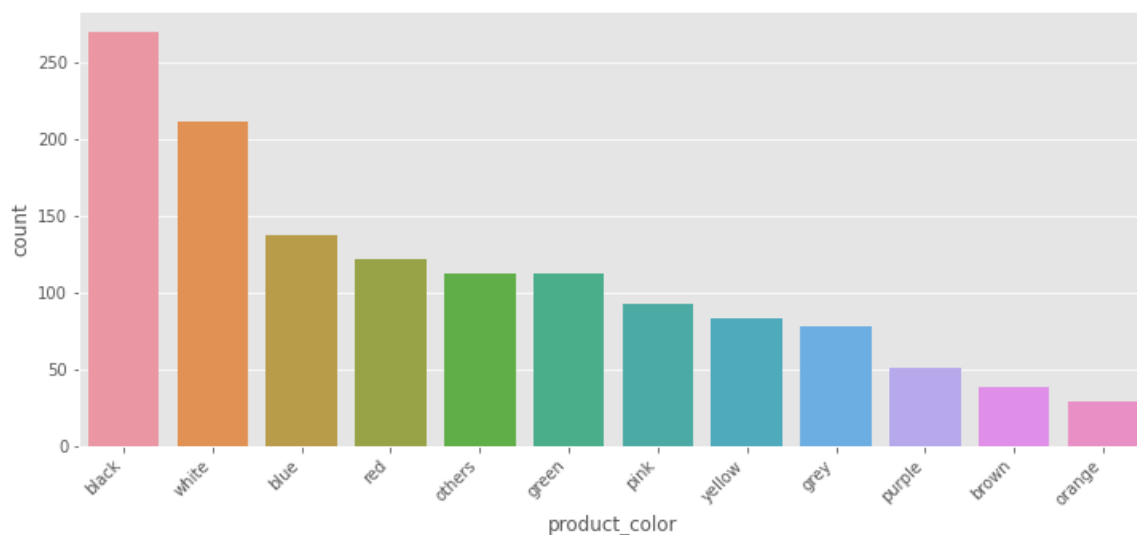
/usr/local/lib/python3.6/dist-packages/seaborn/_decorators.py:43: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



In [56]:

```
# visualize the Product Colours.  
  
fig, ax = plt.subplots(figsize=(12, 5))  
  
sns.countplot(x='product_color', data=df, order=df['product_color'].value_counts()  
( ).index, ax=ax)  
  
ax.set(xlabel='product_color', ylabel='count')  
plt.xticks(rotation=45, ha='right')  
  
plt.show()
```



In [57]:

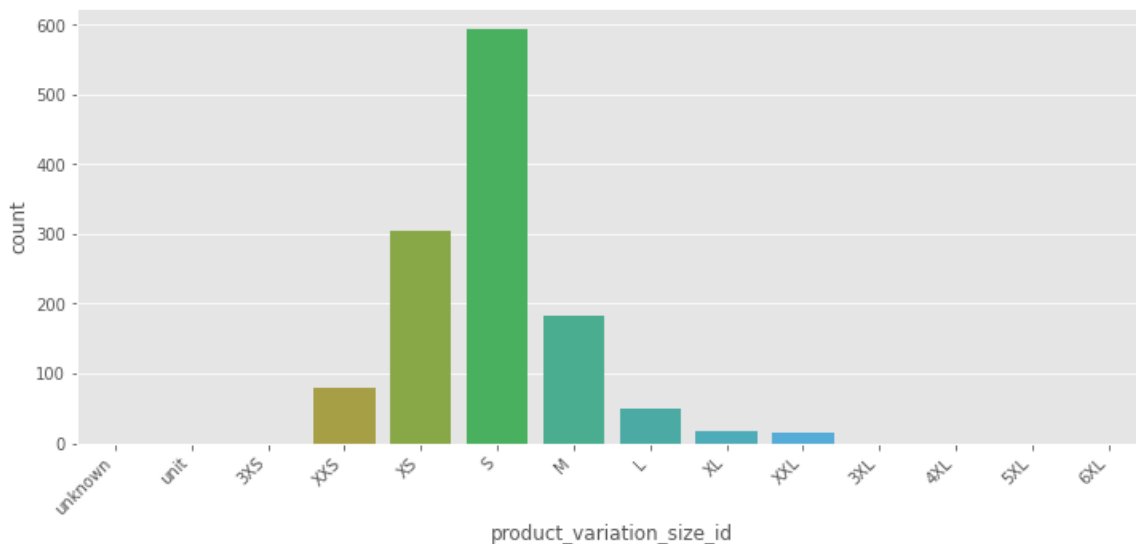
```
# visualize the product_variation_size_id.

fig, ax = plt.subplots(figsize=(12, 5))

size_order = ['unknown', 'unit', '3XS', 'XXS', 'XS', 'S', 'M', 'L', 'XL', 'XXL',
              '3XL', '4XL', '5XL', '6XL']

sns.countplot(x='product_variation_size_id', data=df, order=size_order, ax=ax)

plt.xticks(rotation=45, ha='right')
plt.show()
```



In [58]:

```
# Visualizing product prices (prediction target).

fig = plt.figure(figsize=(12, 5))

fig.add_subplot(2,1,1)
ax1 = sns.distplot(df['listed_price'])
ax1.set_xlim(-10, 260)

fig.add_subplot(2,1,2)
ax2 = sns.distplot(df['retail_price'])
ax2.set_xlim(-10, 260)

plt.tight_layout()
```

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255

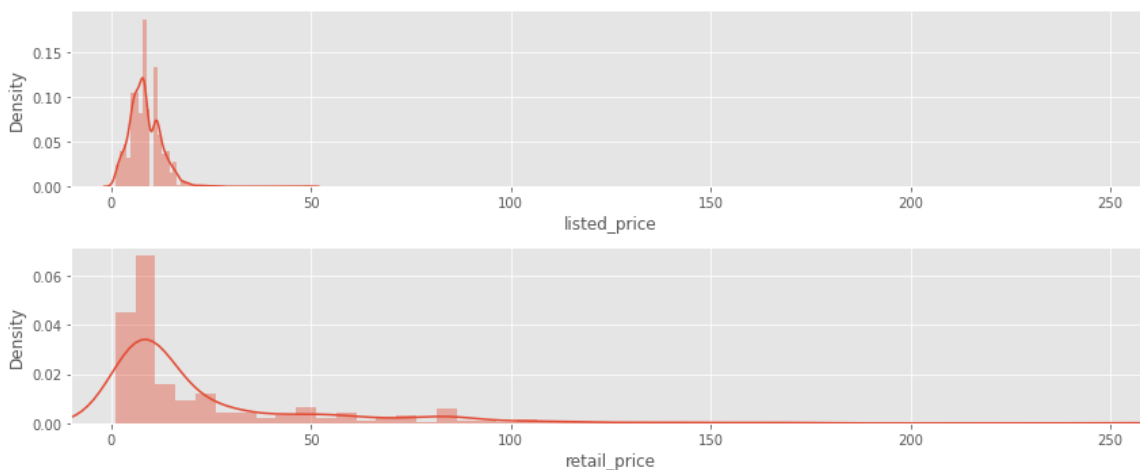
1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255

1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



In [59]:


```
plt.figure(figsize=(12,6))
sb.distplot(df['listed_price'], color='red', label='Listed Price')
sb.distplot(df['retail_price'], color='blue', label='Retail price')
plt.legend();
```

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255
1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255
1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



The plot indicates a right skewed distribution but not very Clear

In [60]:

```
kwargs = {'cumulative':True}
f, axes = plt.subplots(1,2, figsize=(14,6))
f.suptitle('CDF of Listed Price and Retail Price')
sb.distplot(df['listed_price'].values,kde_kws=kwargs, hist_kws=kwargs, color='red', label='Listed Price', ax=axes[0]);
sb.distplot(df['retail_price'].values,kde_kws=kwargs, hist_kws=kwargs, color='blue', label='Retail Price', ax=axes[1]);
axes[0].set(xlabel='Listed Price');
axes[1].set(xlabel='Retail Price');
```

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255

1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/usr/local/lib/python3.6/dist-packages/seaborn/distributions.py:255

1: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



CDFs are more useful in order to visualize the data more efficiently.

CDF of price reveals that 97% of products are listed for less than approximate price 19-20

CDF of Price closely represents the CDF curve of Normal distribution which can be summarized efficiently except for the 3% data

In case of Retail price the distribution is not very much smooth and contains price gaps

In [61]:

```
fig = go.Figure()
fig.add_trace(go.Box(x=df['retail_price'], name='Retail Price'))
fig.add_trace((go.Box(x=df['listed_price'], name='Listed Price')))
fig['layout']['title'] = 'Distribution of Listed Price and Retail Price'
fig.show()
```

With Boxplots we can easily spot the outliers and quartiles

- The Upper fence of Price is at 18 i.e most of the data is priced less than 18
 - There is an item with price of 49 i.e clearly an outlier as it is far away from the Inter Quartile Range (Q3 - Q1)
 - Box plot of Retail price is much more spread out, there is huge difference of 195 between the upper fence and max data point
-

In [62]:

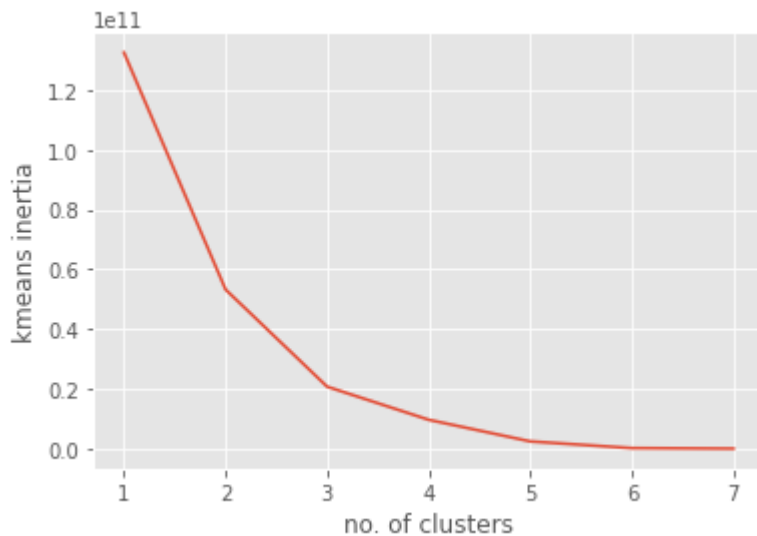
```
#range for units sold  
sorted(df['units_sold'].unique())
```

Out[62]:

```
[1, 2, 3, 6, 7, 8, 10, 50, 100, 1000, 5000, 10000, 20000, 50000, 100000]
```

In [63]:

```
from sklearn.cluster import KMeans  
  
clusters = {}  
for i in range(1,8):  
    kmeans = KMeans(n_clusters=i).fit(df[['units_sold']])  
    clusters[i] = kmeans.inertia_  
  
plt.plot(list(clusters.keys()), list(clusters.values()));  
plt.xlabel('no. of clusters');  
plt.ylabel('kmeans inertia');
```



By performing clustering we can see that units_sold can be clustered in 3 categories (optimal) as the inertia curve smooths out after 3 clusters

In [64]:

```
#order cluster method

def order_cluster(cluster_field_name, target_field_name, df, ascending):
    new_cluster_field_name = 'new_' + cluster_field_name
    df_new = df.groupby(cluster_field_name)[target_field_name].mean().reset_index()
    df_new = df_new.sort_values(by=target_field_name, ascending=ascending).reset_index(drop=True)
    df_new['index'] = df_new.index
    df_final = pd.merge(df, df_new[[cluster_field_name, 'index']], on=cluster_field_name)
    df_final = df_final.drop([cluster_field_name], axis=1)
    df_final = df_final.rename(columns={"index": cluster_field_name})
    return df_final

df['units_sold_cluster'] = KMeans(n_clusters=3).fit(df[['units_sold']]).predict(df[['units_sold']])
df = order_cluster('units_sold_cluster', 'units_sold', df, True)
df.groupby(['units_sold_cluster'])['units_sold'].describe()
```

Out[64]:

	count	mean	std	min	25%	50%	75%	
units_sold_cluster								
0	1056.0	1330.027462	1820.903726	1.0	100.0	1000.0	1000.0	
1	262.0	13778.625954	4857.809727	10000.0	10000.0	10000.0	20000.0	2
2	23.0	63043.478261	22448.887927	50000.0	50000.0	50000.0	75000.0	10

now we have a clear picture of top selling, and price range of products

In [65]:

```
px.scatter(df,x='units_sold',y='rating', color='units_sold_cluster', marginal_y  
='box',title='Rating vs units sold')
```

- Median for rating is 3.85 and the products in top selling cluster has rating between 3.35 to 4.1 seems very reasonable
- Rating is very important to determine the potential of product
- Still there are some products with 5 star rating yet unable to cross the 100-1000 unit sold line
- there are some really bad performing products with rating below 3

In [66]:

```
px.scatter(df,x='rating',y='merchant_rating', color='units_sold_cluster', marginal_y = 'box',title='Merchant Rating vs units sold', opacity=0.7)
```


In [67]:

```
px.scatter(df,x='rating', y='product_variation_inventory', color='units_sold_cluster', title='Product variation vs Rating')
```

In [68]:

```
fig = px.scatter(df,x='rating_count',y='rating', color='units_sold_cluster', title='Rating vs Rating count')
fig.add_trace(go.Scatter(x=np.ones((len(df)))*1103,y=df['rating'],name='Threshold 1'))
fig.add_trace(go.Scatter(x=np.ones((len(df)))*7773, y=df['rating'],name='Threshold 2'))
fig.update_layout(showlegend=False)
```

From above visualization we can conclude that products sold by merchants belonging to cluster 2 and 1 are Top selling,most liked and trusted by buyers

- There's some kind of thresholding that can be done on rating and rating count to separate the 3 categories of products
- still there are few overlapping data points

In [69]:

```
px.scatter(df,x='retail_price', y='listed_price',color='units_sold_cluster',marginal_y='box')
```

Most of the top selling products seems be concentrated to the left where the price difference is much significant

In [70]:

```
px.scatter(df, x='listed_price', y='shipping_option_price', color= 'units_sold_c  
luster', title='Shipping price vs Price')
```

People always prefer paying less shipping charges we can see that most selling products has low shipping charges

The distribution of prices and defining a successful product..

In [71]:

```
print('Median of units sold is',df['units_sold'].median())
print('Mean of units sold is',df['units_sold'].mean())
df['units_sold'].value_counts()
```

Median of units sold is 1000.0

Mean of units sold is 4820.662938105891

Out[71]:

100	396
1000	362
5000	200
10000	163
20000	99
50	50
10	36
50000	17
100000	6
1	3
8	2
7	2
3	2
2	2
6	1

Name: units_sold, dtype: int64

In [72]:

```
def below_ten(units_sold):
    if units_sold < 10:
        return 10
    else:
        return units_sold
```

In [73]:

```
df['units_sold'] = df['units_sold'].apply(below_ten)
df['units_sold'].value_counts()
```

Out[73]:

100	396
1000	362
5000	200
10000	163
20000	99
50	50
10	48
50000	17
100000	6

Name: units_sold, dtype: int64

- The median is 1000 units sold. I will consider for those products with units sold over 1000 units as successful products.
- A quick look at the products with 100000 sales.

In [74]:

```
df[df['units_sold'] == 100000]
```

Out[74]:

	title_translated	listed_price	retail_price	units_sold	uses_ad_boosts	rating	rating_count
1319	2018 New Fashion Women's Tops Sexy Strappy Sle...	5.00	25	100000	1	3.83	1798
1322	Women Stretchy Camisole Spaghetti Strap Long T...	5.77	48	100000	0	4.10	2074
1323	New Aeeival Women Clothing Long Sleeve Autumn ...	8.00	7	100000	1	3.76	1106
1324	Womens Summer Red White and Blue Chiffon Short...	5.00	33	100000	0	3.98	1378
1335	Women Lace Short Sleeve Long Tops Blouse Shirt...	7.00	22	100000	1	3.82	1191
1337	Women's Summer Sexy Sleeveless Turtleneck Mini...	5.67	19	100000	0	3.53	1839

In [75]:


```
# A simple function to create a new label 'is_successful':A simple function to c
reate a new label 'is_successful':

def is_successful(units_sold):
    if units_sold > 1000:
        return 1
    else:
        return 0
```

In [76]:

```
df['is_successful'] = df['units_sold'].apply(is_successful)
#df['is_successful'] = df['units_sold'].apply(is_successful).astype('category')
print('Percent of successful products: ', df['is_successful'].value_counts()[1]
      / len(df['is_successful'])*100)
sb.countplot(data=df, x='is_successful')
plt.show()
```

Percent of successful products: 36.16703952274422



I believe 33% is an appropriate % to be defined as successful sales.

The use of ad boosts to boost the success..

There is an almost equal use of ad boosts by the products.

In [77]:

```
print('Percent of products using ad boosts: ', df['uses_ad_boosts'].value_counts
      ()[1] / len(df['uses_ad_boosts'])*100)
```

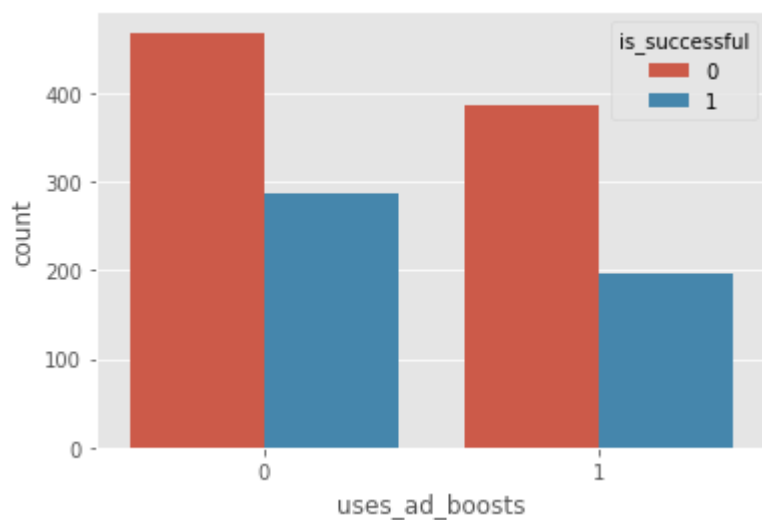
Percent of products using ad boosts: 43.54958985831469

In [78]:

```
sb.countplot(data=df, x='uses_ad_boosts', hue='is_successful')
```

Out[78]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f3c736b96d8>



In [79]:

```
pd.crosstab(df['uses_ad_boosts'], df['is_successful'])
```

Out[79]:

uses_ad_boosts	is_successful	
	0	1
0	469	288
1	387	197

Higher ratings means higher units sold?

In [80]:

```
df['rating']
```

Out[80]:

0 3.76

1 3.57

2 4.03

3 3.10

4 5.00

...

1336 3.91

1337 3.53

1338 4.01

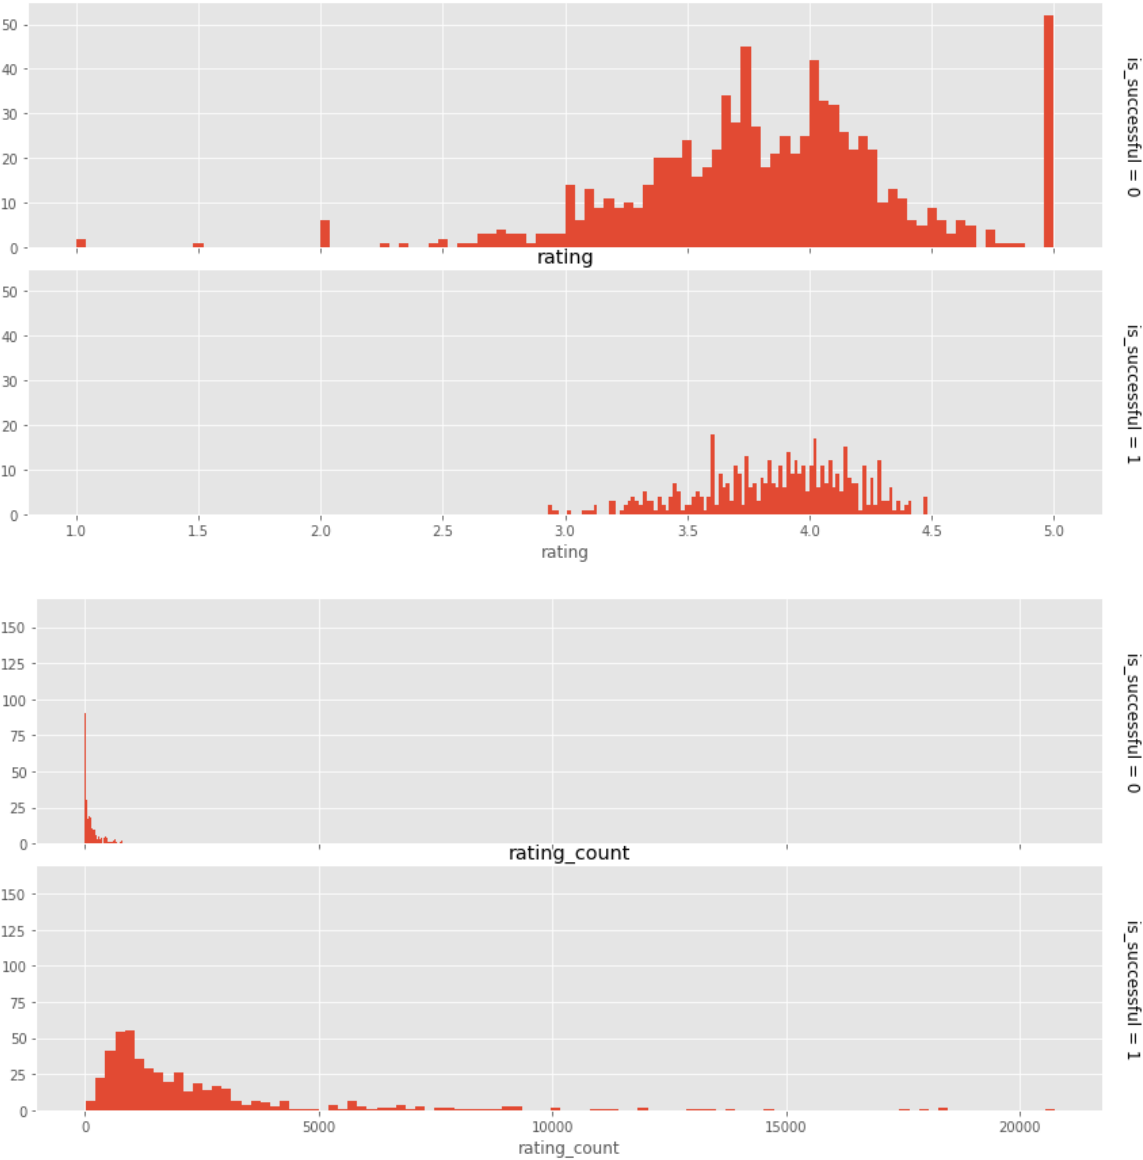
1339 4.01

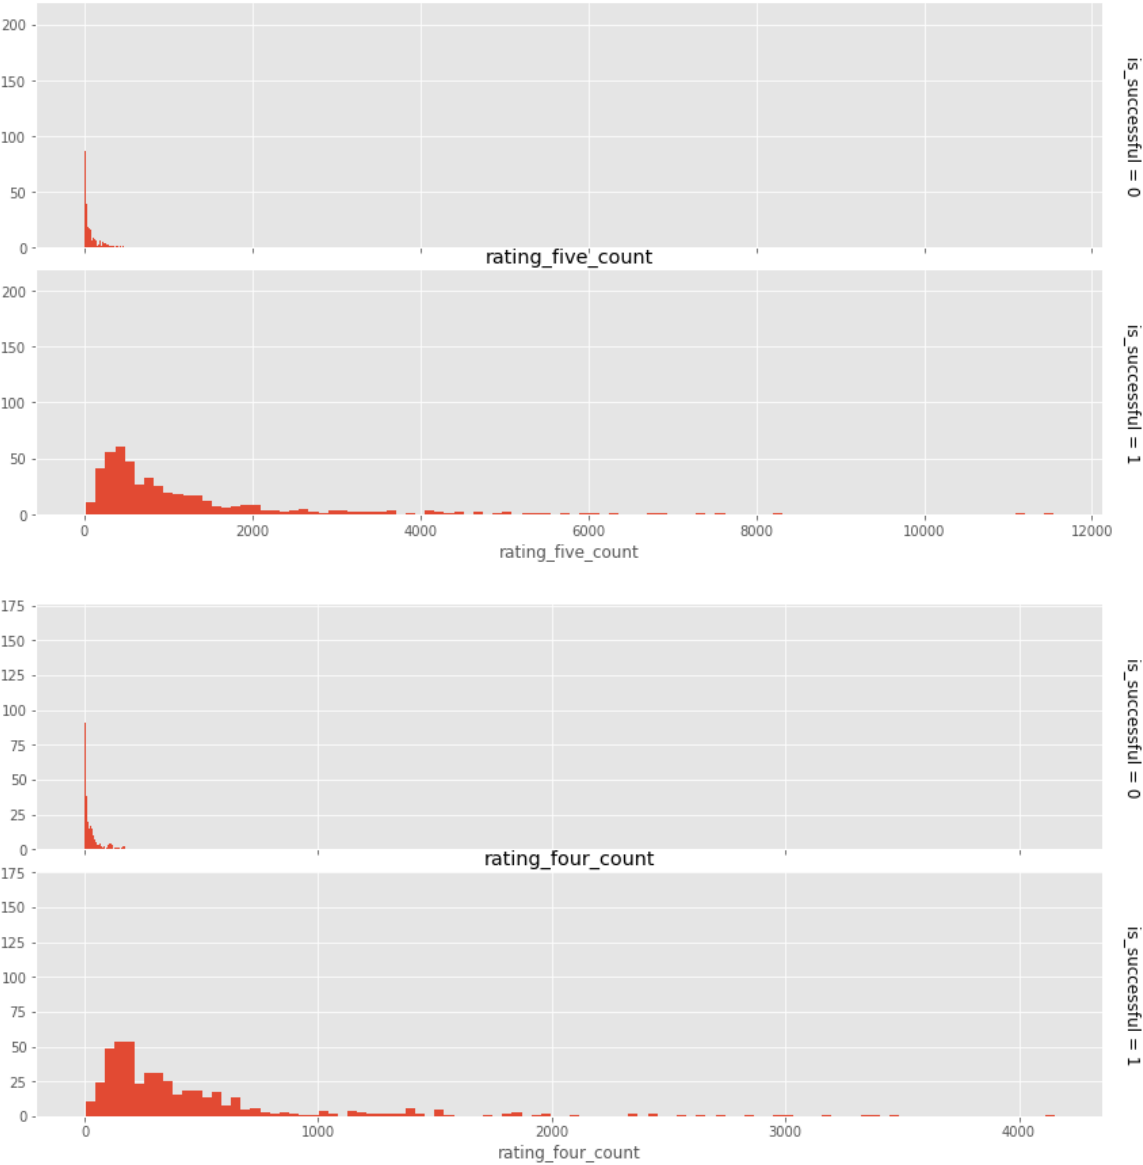
1340 3.60

Name: rating, Length: 1341, dtype: float64

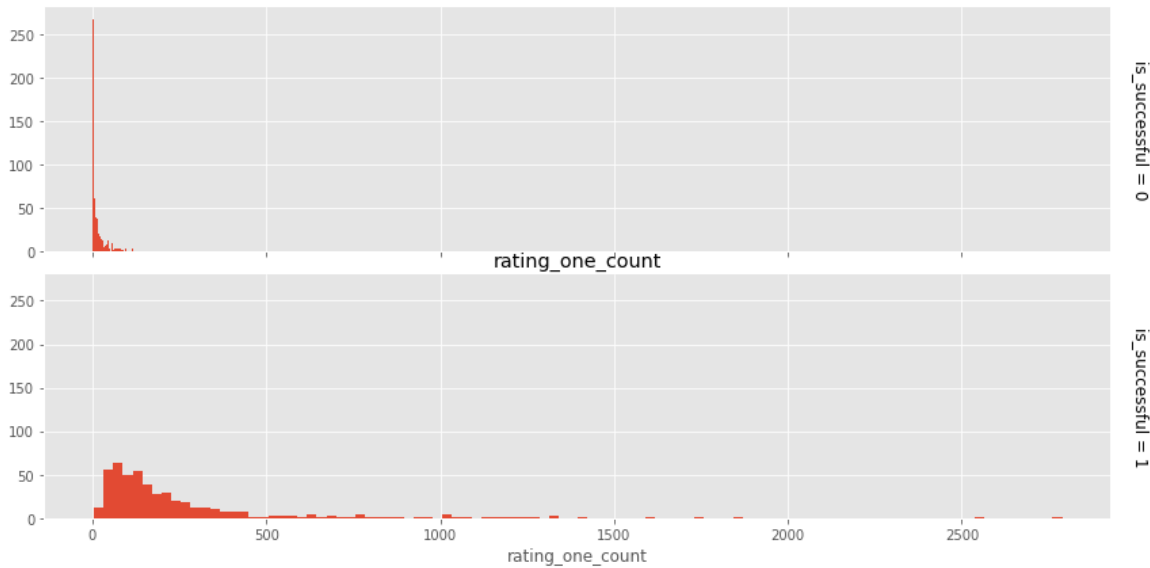
In [81]:

```
ratings_column = ['rating', 'rating_count',  
                  'rating_five_count', 'rating_four_count', 'rating_three_count',  
                  'rating_two_count', 'rating_one_count']  
  
for column in ratings_column:  
    g = sb.FacetGrid(df, row='is_successful', margin_titles=True, aspect=4, height=3)  
    g.map(plt.hist, column, bins=100)  
    plt.title(column)  
    plt.show()
```









In [82]:

```
df.groupby('is_successful').mean()[ratings_column]
```

Out[82]:

	rating	rating_count	rating_five_count	rating_four_count	rating_three_count
is_successful					
0	3.829603	136.151869	65.349299	26.147196	19.109813
1	3.858701	2474.486598	1196.740206	486.369072	365.517526

In [83]:

```
df.groupby('units_sold').mean()[ratings_column]
```

Out[83]:

	rating	rating_count	rating_five_count	rating_four_count	rating_three_count	ra
units_sold						
10	4.388750	1.479167	0.625000	0.229167	0.229167	
50	3.977600	5.920000	3.000000	1.260000	0.720000	
100	3.752626	34.568182	16.002525	6.585859	4.823232	
1000	3.819227	283.121547	136.524862	54.419890	39.781768	
5000	3.818900	882.220000	421.290000	168.720000	127.325000	
10000	3.898282	1847.257669	922.674847	365.361963	258.582822	
20000	3.867576	4507.626263	2161.181818	889.747475	673.959596	
50000	3.903529	10731.941176	5216.764706	2104.882353	1614.058824	
100000	3.836667	15646.833333	7187.166667	3120.500000	2583.500000	

Successful products have more ratings. This is expected as units sold is higher so rating is also higher. However, if I am building a model, I don't think it would be a good idea to keep this column as the huge number of ratings must be the result of higher units sold.

Some conclusions

- The site mainly sells female clothing.
- Higher units sold means higher rating count.
- The use of ad boosts does not seem to have any effect on the units sold and the site may lose revenue from this ads.
- More detailed units sold and inventory levels would have been more helpful for analysis.
- The tags can be improved so that products can be categorised more specifically. This can be done by reducing the number of tags per product, so the merchants are forced to choose their tags more wisely.
- Majority of the products are black and white. This might have been defined wrongly by the merchants. If not the case, the merchants can be encouraged to include more variation to these.